

# Payment Fraud Detection Under Temporal Regime Drift

Severe 1.76% class imbalance compounded by an acute summer behavioral shift across a 2-month test gap.

Extreme Imbalance 1.76% Prevalence

## 1.76%

Only 1 fraud per 57 transactions. Standard cross-entropy heavily under-ranks fraud without surgical threshold and ranking optimization.  
Metric: **Average Precision (PR-AUC)**.

Temporal  
Gap

2-Month Train/Test  
Split

## 60 Days

Training ends July 15; test runs July 16 – September 15. Models relying on static amount thresholds, entity memorization, or short-term memorized patterns decay rapidly.

Empirical  
Audit

Summer Regime  
Shift

## 2x Shift

Diagnostic on labelled folds revealed winter amount-based fraud power **halved** ( $0.41 \rightarrow 0.24$  AP), while short-term device-velocity attacks **doubled** ( $0.095 \rightarrow 0.172$  AP).

**Core Framing:** Static rules and traditional classifiers fail under drift. Solution requires scale-free behavioral baselines and dual temporal adaptation.

0:00 - 0:25

# Strictly-Prior Behavioral Features: 160 Core Signals

Zero target encoding, zero entity IDs, mechanically asserted point-in-time causality.

## 1. Baselines

### Scale-Free Normalization

- `cust_amt_ratio`: Amount / expanding personal mean.
- `cust_amt_robust_z`: Robust MAD-based z-score around personal median:  $\frac{\text{amt} - \text{med}}{\text{MAD} + \epsilon}$ .
- Immune to inflation and currency drift.

## 2. Windows

### Traffic-Matched Ladders

- 10 Trailing Windows: 5m, 15m, 30m, 1h, 3h, 6h, 12h, 24h, 72h, 168h with `closed="left"`.
- Sub-hour windows catch credential-stuffing bursts on high-throughput devices and merchants.

## 3. Acceleration

### Burst Velocity Ratios

- `cust_burst_accel`:  $\frac{\text{cnt}_{1h} \times 24}{\text{cnt}_{24h} + 1}$ .
- Detects sudden transaction spikes relative to the customer's trailing baseline rate.
- High signal for automated fraud attacks.

## 4. Novelty

### Bipartite Expansion

- `is_new_device_for_customer`
- `is_new_location_for_customer`
- `nunique_cust_for_device_prior`
- Tracks account takeover (ATO) and fraudulent device sharing in real time.

**Causality Enforcement:** All features mechanically validated by `assert_strictly_past()`. Pruned 220+ initial candidates down to the Top 160 core features.

0:25 -  
0:55

# Dual-Horizon Tree-Neural Hybrid

Gradient boosted decision trees across dual timeframes fused with continuous deep neural manifolds.

## Trees Dual-Horizon LightGBM Ensemble

- **Full-Train Horizon (Jan 1 – Jul 15):** 882 rounds, num\_leaves=127, lr=0.02. Captures rare long-tail merchants and deep historical categoricals.
- **90-Day Specialist (Apr 16 – Jul 15):** 433 rounds, num\_leaves=63, lr=0.02. Discards stale winter interactions; specializes in summer velocity bursts.
- **Horizon Blend:**  $P_{\text{LGBM}} = 0.48 \cdot P_{\text{Full}} + 0.52 \cdot P_{90\text{d}}$

Local PR-AUC: 0.5284

## Neural Tabular ResNet Regularizer

- **Entity Embeddings:** Learnable continuous representations for high-cardinality categoricals ( $d \approx 1.6 \times \text{card}^{0.56}$ ).
- **3 Residual Blocks:** LayerNorm  $\rightarrow$  Linear(256, 512)  $\rightarrow$  GELU  $\rightarrow$  Dropout(0.15)  $\rightarrow$  Linear(512, 256)  $\rightarrow$  Skip Connection.
- **Tree-Neural Fusion:**  $P_{\text{raw}} = 0.88 \cdot P_{\text{LGBM}} + 0.12 \cdot P_{\text{ResNet}}$

Local PR-AUC: 0.5298 (+0.0014 Lift)

**Why Tree + Neural?** Continuous neural manifolds smooth orthogonal tree step boundaries, rescuing borderline false negatives. Zero pretrained weights used.

0:55 -  
1:30

# Unsupervised Temporal Prediction Propagation

Leave-one-out burst diffusion exploits 3.87x empirical device sibling fraud clustering.

**LOO Diffusion Formula:**  $P_{\text{final}}(i) = (1 - w) \cdot P_{\text{raw}}(i) + w \cdot \frac{1}{2} [\bar{P}_{\text{LOO, cust}}(i) + \bar{P}_{\text{LOO, dev}}(i)]$

$w = 0.50$ ,  $\pm 60\text{m}$  window

Singletons (95.3% of Test)

## Untouched Pass-Through

- **Zero In-Window Sibling Rows:** When a transaction occurs without rapid device or customer siblings,  $\bar{P}_{\text{LOO}} = P_{\text{raw}}$ .
- **Exact Invariance:**  $P_{\text{final}} = P_{\text{raw}}$ . Exactly **95.3% of test rows** pass through completely unmodified.
- Zero false alarms introduced into normal, solitary transactions.

Coordinated Burst (4.7% of Test)

## Qualitative Rescue

**Case Walkthrough: Device X (3 rapid attacks in 18 mins)**

- **Tx 1 (14:02):**  $P_{\text{raw}} = 0.38$  (Borderline False Negative below 0.50 threshold).
- **Tx 2 (14:11):**  $P_{\text{raw}} = 0.74$  (High confidence velocity burst).
- **Tx 3 (14:20):**  $P_{\text{raw}} = 0.64$  (High confidence velocity burst).
- **LOO Sibling Mean for Tx 1:**  $\frac{0.74 + 0.64}{2} = 0.69$
- **Diffused  $P_{\text{final}}$ (Tx 1):**  $0.50(0.38) + 0.50(0.69) = 0.535 \rightarrow$  **Rescued to True Positive!**

**Scientific Rigor:** Touches zero test labels, fits zero test parameters. Elevates burst precision, driving Local PR-AUC  $0.5298 \rightarrow 0.5359 \rightarrow 0.56548$  Public LB.

1:30 -  
2:05

# Systematic Progression & Transparent Failure Modes

Evaluated on chronological holdout (July 1–15, 59,465 txs) with noise floor  $\sigma \approx 0.0020$ .

Ablation

Progressive Milestones

| Iteration       | Key Technical Innovation                            | Local PR-AUC | Public LB |
|-----------------|-----------------------------------------------------|--------------|-----------|
| 0. Baseline     | Raw tabular features, default LightGBM              | 0.1664       | 0.1664    |
| 1. Pipeline     | Expanding baselines, MAD z-scores, windows          | 0.5090       | 0.5105    |
| 2. Calibration  | Removed <code>scale_pos_weight</code> (rank metric) | 0.5204       | 0.5395    |
| 3. Pruning      | Density pruning to Top-160 core features            | 0.5269       | —         |
| 4. Dual-Horizon | 48% Full-Train (882r) + 52% 90-Day (433r)           | 0.5284       | —         |
| 5. Tree-Neural  | 88% Dual-Horizon LGBM + 12% Tabular ResNet          | 0.5298       | —         |
| 6. Diffusion    | LOO Temporal Propagation ( $w = 0.50$ )             | 0.5359       | 0.56548   |

Honest Analysis

Error Modes & Discarded Ideas

- **Primary Failure Mode: Cold-Start Solitary Attacks:**  
First-time fraudsters executing single hits on burner devices with zero prior customer baselines or in-window siblings. Velocity and propagation cannot fire; model relies on static attributes.
- **Discarded Negative Findings:**
  - *Class Weighting* (`scale_pos_weight=55`): Distorted probability calibration under PR-AUC. Removing it was worth **+0.026 LB**.
  - *Forward Features* ( $[t, t + 1h]$ ): Measured +0.024 local lift, but strictly discarded unsubmitted to preserve causal ethics.

# Real-World Production Feasibility

A high-throughput, leakage-free fraud prevention architecture ready for live payment rails.

Latency

Sub-10ms Scoring

## < 10ms

Features update in  $O(1)$  time via incremental in-memory state. Tree ensemble and PyTorch ResNet forward pass complete under 5ms on standard CPU.

Governance

100% From Scratch

## Zero Pretrained

Zero foundation models, zero external pretrained weights, zero third-party licensing risks. Fully reproducible using standard `lightgbm` and `torch`.

Streaming

In-Memory Ring Buffer

## Redis Buffer

LOO burst propagation implements naturally in live payment gateways using a 60-minute in-memory Redis sliding buffer, flagging coordinated attacks asynchronously.

**Final Takeaway:** Causally sound scale-free features + dual temporal horizons + unsupervised burst diffusion transforms temporal drift from an adversary into our greatest advantage.

2:40 -  
3:00