

REACT 2026 Datathon — Methodology & Compliance Disclosure

Team Overfit & Overcaffeinated

2026-09-08

Methodology disclosure — prediction propagation

For the reproducibility notebook. Written to be read by a competition reviewer, and to disclose the one part of our pipeline whose compliance is a judgement call rather than a plain fact.

The rule

“Every engineered feature for a transaction at time t may only use information strictly before t .”

and its clarification:

“Feature computation that uses only the raw, non-target columns of test.csv in a strictly-past-only way is fine (e.g., a device’s known transaction history can include test-period rows that occurred earlier than the row being scored); using the fraud label anywhere near test.csv is not, since it does not exist for you.”

What our submission does

Every engineered feature is strictly prior. All ~197 features — entity expanding aggregates, trailing windows, novelty and fan-out counts, frequency shares, calendar terms — are computed so that a row at time t sees only rows strictly before t . This is enforced mechanically, not by inspection: `leakage_assertions()` checks that first-occurrence rows carry no prior statistics, and `scripts/test_compliance.py` verifies empirically that perturbing a later row cannot change any earlier row’s feature values.

No label is touched anywhere near test.csv. We use no target encoding, no per-entity fraud rates, and no per-subgroup label aggregation at any grain finer than the whole training set. Raw `customer_id` / `merchant_id` / `device_id` / `transaction_id` never enter the feature matrix in any encoding; they are join keys only, asserted at every `prepare_lgb_frame()` call.

One post-inference step reads adjacent rows, and it is the reason for this document. After the model has trained and predicted, we apply an entity propagation blend:

```
final(i) = 0.5 * raw(i) + 0.5 * mean( raw(j) : j shares i's customer or device
                                         and |t_j - t_i| <= 30 minutes, j != i )
```

The window is symmetric, so for a row at time t it includes rows up to $t + 30$ minutes.

Why we believe it is within the rules

- It is **not an engineered feature**. The model never sees it. It is arithmetic over model outputs, applied after training and inference are complete, and the rule as written governs feature computation.
- It **uses no labels of any kind**. Its only inputs are the model’s own predictions, entity identifiers, and timestamps. Nothing is fitted to `test.csv`; there is no parameter estimated from test data. The window width and blend weight were selected on labelled validation windows drawn from `train.csv`.
- It encodes a **behavioural pattern, not entity targeting**: fraudulent transactions cluster in time within an entity. On training data, $P(\text{a sibling transaction is fraud} \mid \text{this one is fraud})$ is 3.87x the base rate for devices and 2.63x for customers. The blend expresses that clustering, which the rules explicitly invite (“do groups of customers, devices, and merchants form suspicious clusters that no single transaction reveals on its own?”). It does not key on which entity it is.

Why we are disclosing it rather than leaving it implicit

The rule’s evident purpose is temporal causality, and the clarifying example permits test-period rows only where they “occurred earlier than the row being scored”. A symmetric window does not meet that description. We think the narrow reading is correct — the sentence governs *features* — but we do not think a reviewer should have to discover the question on their own, so we are stating it plainly and accepting the organiser’s judgement.

Scope of the effect. The blend changes **4.7% of test rows** (3.6% via customer, 4.1% via device); the remaining 95.3% have no in-window neighbour and pass through unchanged. It is a refinement of the ranking, not the basis of it.

Our compliant alternative

Because this is a judgement call, we also prepared a submission with the propagation step removed and nothing else changed — every feature strictly prior, no post-inference blending, no future information anywhere. If the organiser reads the rule more strictly than we do, that submission stands on its own and we would ask for it to be scored instead.

Our two selected submissions: - **submission.csv** (public LB 0.56548) – includes the propagation step described above. Reproduced by `REACT_2026_champion_reproducibility.ipynb`. - **CANDIDATE_v9c_compliant_0.5280.csv** (public LB 0.55368) – strictly-past only, no propagation, nothing to argue. Reproduced by `REACT_2026_v9c_reproducibility.ipynb`.

Both notebooks are in this folder alongside this document.

What we excluded on the same reasoning

We built and measured forward-looking window features (counts and sums over $[t, t+w]$, time-to-next-transaction). They were worth roughly +0.024 local PR-AUC — our largest single measured gain. **We discarded them and did not submit them**, because unlike the propagation step these are unambiguously engineered features reading information after t , which the rule forbids. They are absent from every submitted model, and `assert_strictly_past()` now fails the build if any such column reaches a feature matrix.